Enterprise BI for CDA Governance

A working guide to the business-intelligence practices that transfer to communications data acquisition — what to know, what to be able to talk around, and what not to claim

Prepared for Michael Rossi
October 2026

Contents

Contents 2
1. The data platform and its layers 4
2. Data modelling: facts, dimensions and keys 5
3. Slowly changing dimensions and point-in-time correctness 6
4. Master data management and identity resolution 7
5. ETL/ELT and the pipeline 8
6. Data quality 9
7. The data catalogue and classification 10
8. Lineage and provenance 11
9. Access control in data platforms 12
10. Retention, archival and deletion 13
11. Data governance as an organisational function 14
12. Platform operations and capacity 15
13. Vocabulary cheat-sheet 16
14. Conversation scripts 17

Who this is for and how to use it

This guide assumes you have designed and governed large systems but have not run a BI or data-warehouse function yourself. It teaches the practices from the ground up — enough to understand what a BI professional means, to recognise the same practice when it appears in a CDA design under a different name, and to hold a credible conversation with either side. Each chapter has the same shape: what the practice is, the vocabulary, how it shows up in CDA, what you can say about it from an architect's standpoint, and what you should not claim.

The honest position to take throughout: *I have governed the systems these practices run inside, I understand them as an architect understands them, and I have not personally built a warehouse.* That is a strong position. The disciplines below are the ones a CDA design authority has to hold people to; the person holding them to account does not need to have written the ETL.

Chapter	Practice	CDA name for it
1	The data platform and its layers	The six-layer reference architecture
2	Data modelling: facts, dimensions and keys	Entity, events, identifiers
3	Slowly changing dimensions and point-in-time correctness	Interval bindings; time-segmented resolution
4	Master data management and identity resolution	Identifier-to-subscriber resolution
5	ETL/ELT and the pipeline	Ingest and normalisation
6	Data quality	Quality controls, volume reconciliation
7	The data catalogue and classification	Field-level catalogue
8	Lineage and provenance	Provenance statement; evidential weight
9	Access control in data platforms	Scope enforcement
10	Retention, archival and deletion	Retention period; destruction; crypto-shredding
11	Data governance as an organisational function	The design authority
12	Platform operations and capacity	Capacity planning; feed monitoring
13	Vocabulary cheat-sheet	
14	Conversation scripts	

1. The data platform and its layers

What it is

Enterprise BI is the discipline of taking data out of the systems that run the business (operational or transactional systems — the CRM, the billing platform, the order system) and putting it somewhere it can be queried for understanding rather than for operation. The 'somewhere' has had several names over thirty years — data warehouse, data mart, data lake, lakehouse — but the structure is constant.

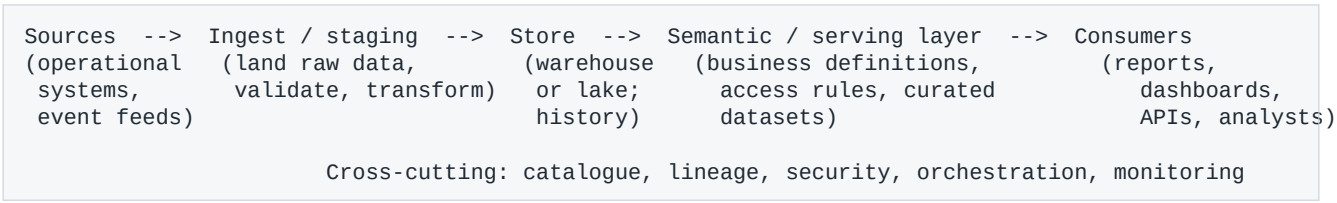

Term	Meaning
Operational / transactional system (OLTP)	A system that runs the business and is optimised for many small writes. The source. In CDA: the IMS nodes, HSS, PGW, NAT, CRM
Data warehouse	A store optimised for querying large volumes of historical data, structured into a model. Kimball-style (star schemas) or Inmon-style (normalised then marts). In CDA: the retention store
Data lake	Raw data stored cheaply in its native format, structured on read. Lakehouse: a lake with warehouse-like tables and governance on top
Data mart	A subject-specific slice of the warehouse for one consumer group
Staging / landing zone	Where raw extracts land before transformation. In CDA: the ingest layer, where truncation and classification happen
Semantic layer	The layer that defines business terms over physical tables — 'revenue' means these columns combined this way — and applies access rules. In CDA: disclosure management, where scope and caveats are applied
Serving layer / consumption	Reports, dashboards, APIs, extracts. In CDA: the external interface to the acquiring authority
Medallion architecture (bronze / silver / gold)	A lakehouse convention: raw, cleaned, curated. Maps to L1–L2 (bronze), L3 (silver), L5 (gold)

What you can say

> *'A CDA disclosure capability is a data platform with a legally defined consumer. The layers are the ones any warehouse architect would draw — sources, staging, store, a semantic layer that applies the rules, a serving interface — and the controls sit on the boundaries between them. What is different is that the semantic layer's rules come from an authorisation rather than a business definition.'*

What not to claim

Do not claim to have chosen between Kimball and Inmon, tuned a warehouse, or built a lakehouse. Say you have governed platforms that had one, and that the layer boundaries are where you have always put the controls.

2. Data modelling: facts, dimensions and keys

What it is

Warehouse data is modelled as **facts** — things that happened, with measures (a sale, a call, a login) — and **dimensions** — the things facts are about (customer, product, date, location). A star schema is one fact table surrounded by dimension tables, joined on keys. The fact table is large and append-only; dimensions are smaller and describe.

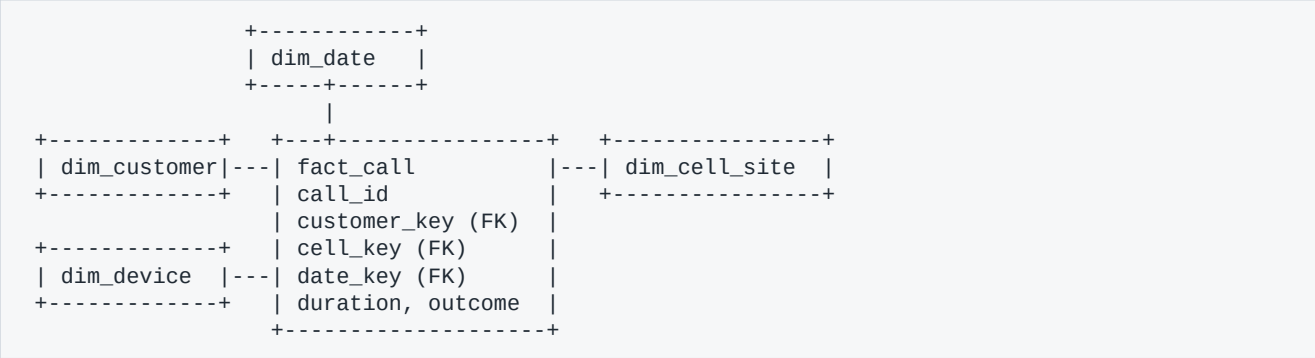

Term	Meaning	CDA counterpart
Fact	An event or transaction with measures	Events data: a session, a registration, a message
Dimension	Descriptive context a fact refers to	Entity data: the subscriber, the device, the cell site
Grain	What one row of a fact table represents — one call, one call leg, one bearer event	The design authority's decision about what one retained record is; the ICID correlation question is a grain question
Business key / natural key	The identifier the source uses (MSISDN, account number)	MSISDN, IMSI, IMPU — reusable, so not safe as a sole key
Surrogate key	A warehouse-generated key that is stable even when the business key is reassigned	An internal subscription identifier that survives number recycling
Degenerate dimension	A transaction identifier stored on the fact itself	ICID, charging ID
Conformed dimension	A dimension shared, with the same meaning, across fact tables	One cell-site reference used by both 4G and 5G records; one subscriber reference used by core and IMS records

What you can say

> *'The entity/events split in the Act is the dimension/fact split in a warehouse, with a legal threshold attached. The thing a modeller would ask first — what is the grain of a retained record? — is the question that decides whether one call is one row or six, and therefore whether correlation is done at ingest or at query time.'*

What not to claim

Do not claim to have designed star schemas. You can say you have reviewed data models for large systems and that grain and key stability are the two things you always check.

3. Slowly changing dimensions and point-in-time correctness

What it is

Dimensions change: a customer moves, a product is renamed, a number is reassigned. A warehouse has to decide whether to overwrite (and lose history) or keep history. The standard answer is the **Type 2 slowly changing dimension**: each version of a dimension row has a valid-from and valid-to date, and a fact joins to the version that was in force when the fact happened.

```
Type 1 (overwrite)            Type 2 (versioned)
msisdn   customer             msisdn   customer   valid_from   valid_to     current
0771...  B                    0771...  A          2025-06-01   2026-02-14   N
                              0771...  (none)     2026-02-15   2026-02-19   N
                              0771...  B          2026-02-20   9999-12-31   Y

Type 1 answers 'who has this number?'  Type 2 answers 'who had it on 10 Feb?'
```

The classic failure is a **Type 1 join against historical facts**: every call ever made on the number is now attributed to customer B, including the ones customer A made. In a sales warehouse that misstates a region's revenue. In CDA it attributes someone else's communications to the wrong person.

Term	Meaning
SCD Type 0 / 1 / 2 / 3 / 6	Never change / overwrite / version with dates / keep previous value in a column / hybrid
Effective dating, bitemporal	Valid time (when it was true in the world) versus transaction time (when the system learned it). Bitemporal stores both — relevant when a binding is corrected after the fact
Point-in-time (PIT) query / 'as-at'	A query that joins facts to the dimension version in force at the fact's time
Late-arriving dimension	A fact arrives before its dimension row exists — the subscriber is not yet in the reference data
Gap	A period with no valid version — the number was unallocated

What you can say

> *'Every identifier lookup is only valid with respect to a time — that is a Type 2 dimension with explicit gaps. The question I ask of any resolution design is whether it is bitemporal: does it record when the binding was true and when we learned it, so that a correction after the fact is auditable rather than silent?'*

What not to claim

Do not claim to have implemented SCDs. Say you have seen what happens when a system joins current reference data to historical records, and that you treat as-at correctness as a non-negotiable requirement.

4. Master data management and identity resolution

What it is

Master data management (MDM) is the practice of maintaining one authoritative version of the core entities an organisation cares about — customer, product, supplier, location — when those entities appear, inconsistently, in many systems. The output is a **golden record** per entity with a **cross-reference** to each source system's key. The hard part is **entity resolution**: deciding that 'M. Rossi, 12 High St' in the CRM and 'Michael Rossi, 12 High Street' in billing are the same person.

Term	Meaning	CDA counterpart
Golden record	The single trusted version of an entity	The subscription record in HSS/UDM plus the CRM account
Cross-reference (xref)	The table mapping each source key to the golden record	The binding store: MSISDN, IMSI, IMEI, IMPU, account — each with validity
Deterministic matching	Keys match exactly by rule	The only kind allowed in the CDA resolution path
Probabilistic / fuzzy matching	A score from similarity of name, address, date of birth; above a threshold, merge	Ruled out: nobody can explain a score in a witness box
Survivorship rules	Which source wins when they disagree	System-of-record nomination per attribute
Hierarchy / relationship management	Households, corporate families	Identifier-to-identifier hops — each one a potential disclosure about a different person
Stewardship	Humans resolving the cases rules cannot	The human review point; the 'insufficient information' nil return

What you can say

> *'Identity resolution in CDA is an MDM cross-reference with three constraints MDM does not usually have: every binding is time-qualified, matching must be deterministic and explicable, and every hop from one identifier to another is a disclosure in its own right that needs its own authority. If a vendor's resolution engine has a match score in it, it is in the wrong place.'*

What not to claim

Do not claim to have run an MDM programme. You can say you have governed systems where the same party appeared under several keys and that you know where the probabilistic shortcuts tend to be hidden.

5. ETL/ELT and the pipeline

What it is

ETL — extract, transform, load — is the pipeline that moves data from sources to the store. ELT loads first and transforms inside the store, which is the modern default on cloud platforms. Pipelines are batch (nightly extracts) or streaming (events as they happen) and are run by an **orchestrator** that schedules jobs, handles dependencies and retries.

Term	Meaning	CDA counterpart
Extract / connector / CDC	Reading from the source; change data capture reads the source's own change log rather than querying it	Feeds from IMS nodes, NAT, HSS — some pushed by the node, some pulled
Transform	Parsing, cleansing, canonicalising, deriving, joining	Normalisation, truncation of content-bearing fields, classification, correlation on ICID
Load	Writing to the store, typically into partitions	Writing to time-partitioned, classified, immutable partitions
Idempotency	Re-running a job does not double-load	A replayed feed must not create duplicate retained records
Backfill / replay	Loading historical data after a gap or a logic change	Re-ingesting after a feed outage — with the audit showing the replay
Schema drift / schema evolution	The source changes shape; the pipeline must notice	A node upgrade adds or renames a field — the capability-impact question
Quarantine / dead-letter queue	Records that fail validation are set aside, not dropped silently	Rejected records with reasons; the completeness statement counts them
Orchestrator	Airflow, Azure Data Factory, dbt Cloud, Dagster	Whatever schedules and monitors the ingest jobs
Streaming	Kafka, Event Hubs, Kinesis; process records as they arrive	Near-real-time ingest for the urgent path

What you can say

> *'The ingest pipeline is where the irreversible decisions happen: what gets truncated, what gets classified, what gets rejected. In a warehouse you can usually reload from source if you got the transform wrong. Here the source has deleted its copy on PECR grounds, so the pipeline has to be right first time and its rejects have to be counted, not discarded.'*

What not to claim

Do not claim to have written pipelines in a particular tool. You can speak about integration platforms you have designed — message flows between enterprise systems, reconciliation, idempotency, replay — which is the same discipline.

6. Data quality

What it is

Data quality management is the practice of defining what 'good' means for a dataset, measuring it continuously, and acting on failures. The DAMA DMBOK dimensions are the common vocabulary: completeness, accuracy, timeliness, consistency, validity, uniqueness. Quality is enforced by **rules** (checks run in the pipeline), **reconciliation** (comparing counts and totals between source and target), and **monitoring** (dashboards and alerts on the rules).

Practice	What it looks like in BI	In CDA
Source-to-target reconciliation	Row counts and control totals from the extract compared to what loaded; a nightly report; a threshold that pages someone	Volume reconciliation against source-side counts per feed per period — the course's single most important operational control
Data quality rules	Not-null, referential integrity, domain values, ranges; implemented in Great Expectations, dbt tests, Soda, or the platform's own framework	Identifier format validation, mandatory field checks, classification completeness, correlation success rate
Freshness / SLA monitoring	'The sales table is 6 hours stale' alerts	Generation-to-availability latency reported at the tail (p99), not the mean
Anomaly detection on volumes	Today's rows versus the expected range for this weekday	The feed that silently stopped; the IoT population that doubled
Data quality scorecards	Per dataset, per dimension, trended	The assurance pack's quality section
Issue management	A data quality issue log with owner and remediation	The error taxonomy with detection, notification and enumeration of affected disclosures

> **TALKING POINT**
>
> The one sentence to remember: **measure against independent expectation, never against the feed's own count.** A feed that reports 'I sent 0 records' agrees perfectly with a store that received 0. Expectation comes from history, from the source node's own counters, or from a parallel path.

What you can say

> *'Every BI platform I have seen fail did so quietly: a feed stopped, the counts matched because both sides were zero, and nobody noticed until the month-end report looked wrong. In CDA there is no month-end report. The reconciliation has to be designed in from day one and it has to compare against something the pipeline itself does not produce.'*

7. The data catalogue and classification

What it is

A data catalogue is the inventory of what data exists, what it means, where it came from, who owns it, how sensitive it is and who may use it. Modern catalogues (Collibra, Alation, Microsoft Purview, Atlan, open-source DataHub) crawl the platform, attach business definitions, and record classification tags such as PII, confidential, financial. Classification drives access policy — ideally automatically, often only on paper.

Term	Meaning	CDA counterpart
Business glossary	Agreed definitions of business terms	The interpretation record: what 'data identifying the sender' means in this operator
Technical metadata	Tables, columns, types, sources	The field-level catalogue
Classification / sensitivity tags	PII, special category, confidential	Entity / events / content; ICR flag; third-party-data flag
Data owner / steward	Accountable person / day-to-day custodian per dataset	Design authority sign-off per arguable classification
Policy-as-code / tag-based access	Access rules evaluated from the tags at query time	Classification-driven scope enforcement at the gateway
Catalogue drift	The catalogue says one thing, the schema another	The classification spreadsheet no code references — the course's 'compliance by documentation'

What you can say

> *'The test of a catalogue is whether anything reads it. If the pipeline and the access layer consume the classification, it is a control; if only the governance team does, it is a document. In CDA the classification has legal effect, so it has to be the former.'*

8. Lineage and provenance

What it is

Lineage records how data moved and changed from source to output: which source tables, which transformations, which job runs, which version of the logic. Table-level lineage is common; column-level lineage is harder and more valuable. It is used for impact analysis (what breaks if this source changes?), for debugging (why is this number wrong?), and for audit (prove where this figure came from).

Term	Meaning	CDA counterpart
Upstream / downstream	Sources feeding a dataset / consumers of it	Impact analysis for the capability-impact question
Column-level lineage	This output column derives from these input columns via this logic	The provenance statement per field
Run / job lineage	Which execution produced this data	Pipeline version and ingest time on every record
Reproducibility	Re-running the same logic on the same inputs gives the same output	Repeatability as a component of evidential weight
OpenLineage, Marquez	Open standards and tools for capturing lineage	

What you can say

> *'Provenance for a court is column-level lineage with a run identifier: which node, which feed, which pipeline version, when. If the platform cannot answer that for a figure on a dashboard, it cannot answer it for a disclosure, and the evidential case rests on assertion.'*

9. Access control in data platforms

What it is

Data platforms control access at several levels: to the platform, to a dataset, to rows within it (**row-level security**: a regional manager sees only their region), and to columns (**column-level security** or **dynamic data masking**: analysts see a hashed national insurance number). Policies are role-based (RBAC) or attribute-based (ABAC: rules on user attributes and data tags evaluated at query time).

Term	Meaning	CDA counterpart
RBAC	Permissions by role	Insufficient alone: a valid role with no live authorisation must be able to do nothing
ABAC / policy-based access	Rules evaluated per request from user, resource and context attributes	Scope enforcement: authorisation type, data classes, period and subject intersected with the request, per query
Row-level security	Filter rows by policy	Period and subject restriction
Column-level security / masking	Hide or transform columns by policy	Field-set restriction by classification; truncation
Separation of duties	Nobody both builds and approves; admins cannot read data	Compartmentation of provisioning, data, disclosure and audit
Break-glass	Emergency elevated access, logged and reviewed	Vendor and urgent-path access
Query audit	Every query logged with user, time, text, rows returned	The requested / authorised / returned triple per disclosure

What you can say

> *'Role-based access tells you who may use the system; it does not tell you whether this query is lawful. That needs the authorisation as an input to the policy decision — attribute-based control evaluated per request. Most platforms can do it; few programmes configure it that way, because it is easier to trust the operator.'*

10. Retention, archival and deletion

What it is

Every regulated data platform has a records retention schedule (how long each class of data is kept and why), tiered storage (hot, warm, cold, archive), and a deletion process. Since 2018 the right to erasure has forced platforms to be able to find and delete an individual's data, including in derived datasets and backups — which most could not do, and many still cannot.

Term	Meaning	CDA counterpart
Retention schedule	Per-class retention period and legal basis	The notice period per data family
Partition pruning / lifecycle policy	Drop whole partitions when they age out	Time partitioning as the deletion mechanism
Legal hold	Suspend deletion for specified data under litigation or investigation	A designed exception with its own audit
Soft delete / tombstone	Mark as deleted but keep the row	Not acceptable as deletion of retained data
Crypto-shredding	Encrypt per partition or per subject; destroy the key to delete	The answer to backups outliving the retention period
Derived data	Aggregates, extracts, ML features built from the source	Correlated copies, indexes, analytics extracts — all are copies for deletion purposes
Deletion evidence / certificate	Proof that deletion happened	Retained longer than the data it describes

What you can say

> *'The question I ask every platform is: when you delete, what happens in the backups? The answer is usually a pause. The only robust answer is that the backups never held readable data in the first place — per-partition keys, destroyed on expiry.'*

11. Data governance as an organisational function

What it is

Data governance is the people-and-process layer: who is accountable for which data, how decisions about it are taken and recorded, how policy is set and enforced. Typical structures: a data governance council or board, domain data owners, data stewards, a data protection officer, and an architecture review board. Frameworks: DAMA DMBOK, DCAM, and in government the Data Ethics Framework and the Data Maturity Assessment.

Role	BI meaning	CDA counterpart
Data owner	Accountable executive for a data domain	The controller-level accountable owner for the retained store
Data steward	Day-to-day custodian: definitions, quality, access requests	The person who maintains the field catalogue and the interpretation record
Data governance council	Cross-functional decision body for data policy	The design authority with a standing decision list
Architecture review board	Reviews and approves designs against standards	The design authority in its review function
DPO	Statutory data protection adviser	Consulted on the DPIA; part of the governance
Data policy / standards	Documented rules for classification, retention, quality, access	The negative-decision register, classification rules, nil-return vocabulary
Decision log / ADRs	Architecture decision records: context, decision, consequences, date	The decision list the inspector will ask for

What you can say

> *'Governance that only reviews documents is a reading group. The design authority I run keeps a decision list — classification of arguable fields, truncation point, scope mechanism, hop bound, urgent-path scope — with the rationale and date against each, and the test is whether any of them can be produced two years later.'*

This is the chapter where your actual experience is strongest; architecture governance on large programmes is the same practice. Say so.

12. Platform operations and capacity

What it is

Running a data platform means monitoring pipelines and query performance, managing storage growth, controlling cloud cost, and handling incidents. **DataOps** borrows DevOps practice: version-controlled transformations, automated tests, CI/CD for pipelines, observability.

Practice	Meaning	CDA counterpart
Observability / data observability	Monitoring freshness, volume, schema, distribution and lineage automatically (Monte Carlo, Bigeye, open-source Elementary)	Feed monitoring; expected-versus-actual volumes; schema-drift detection
Capacity planning	Forecasting storage and compute from growth in sources, consumers and retention	NAT translation records dominate; indexes and replicas; steady state at twelve months
Cost management (FinOps)	Attributing and controlling cloud spend	Cost-recovery claims under the IPA depend on being able to attribute cost to the capability
CI/CD for data	Pipelines and models tested and deployed through automation	Change control with the capability-impact question in the pipeline
Environment separation	Dev / test / prod with synthetic or masked data in non-prod	Test data must never be real retained data; rehearsal of the assurance pack on synthetic data
Incident management	Runbooks, on-call, post-incident review	The error taxonomy; three-regulator notification routing

What you can say

> *'The capacity question nobody asks is the ratio of the biggest low-value dataset to the one everyone sized for. In a web platform it is the clickstream; in CDA it is NAT translation logs. I ask for that ratio, and whether it was measured or assumed.'*

13. Vocabulary cheat-sheet

Terms you should recognise instantly and be able to use correctly, with the one-line definition. Marked (*) where the CDA course uses a different word for the same thing.

Term	One line
OLTP / OLAP	Transactional systems that run the business / analytical systems that query it
Star schema, snowflake schema	Fact table with denormalised dimensions / with normalised dimensions
Grain	What one fact row represents
Fact, dimension, measure	Event / context / numeric value on the event (* events data, entity data)
Surrogate key, natural key	Warehouse-generated stable key / source identifier
SCD Type 2	Versioned dimension rows with valid-from/to (* interval bindings)
As-at / point-in-time query	Join to the version in force at the fact's time (* time-segmented resolution)
MDM, golden record, xref	Master data; the trusted entity record; the source-key mapping (* binding store)
Entity resolution, deterministic vs probabilistic	Deciding two records are the same thing; by rule vs by score
ETL / ELT, CDC	Pipeline patterns; change data capture from source logs (* ingest and normalisation)
Idempotent, replay, backfill	Safe to re-run; re-processing after a gap
Staging, bronze/silver/gold	Landing zone; raw/cleaned/curated tiers
Data quality dimensions	Completeness, accuracy, timeliness, consistency, validity, uniqueness
Reconciliation, control totals	Comparing counts/sums between source and target (* volume reconciliation)
Freshness, SLA, p99	How stale; the commitment; the tail latency
Catalogue, glossary, classification tag	Inventory; definitions; sensitivity label (* field-level catalogue)
Lineage (table / column level)	Where data came from and how it changed (* provenance)
RBAC / ABAC, RLS / CLS, masking	Role / attribute access; row / column security; transforming sensitive columns (* scope enforcement)
Separation of duties, break-glass	No one both does and approves; emergency access
Retention schedule, lifecycle policy, legal hold	How long and why; automatic ageing; suspend deletion
Crypto-shredding	Delete by destroying keys
Data owner, steward, governance council, ARB	Accountable exec; custodian; decision body; design review (* design authority)
DAMA DMBOK	The standard body of knowledge for data management
DataOps, data observability	DevOps for data; automated monitoring of pipelines and datasets
Semantic layer, certified dataset	Business definitions and rules over physical data; a governed output (* disclosure management)
Data contract	A versioned agreement between a producer and consumer of a dataset on schema, semantics and SLA (* data supply agreement)
Data mesh, data product	Domain-owned datasets with product thinking; contrast with central warehouse

14. Conversation scripts

Short exchanges for the situations where the BI parallel is likely to come up. Each gives the prompt you might hear, a response that is accurate about your experience, and the follow-up question that moves the conversation onto ground you control.

You hear	You say	Then ask
'Have you worked with data warehouses before?'	'I have governed the platforms they sat in and reviewed the models and pipelines, rather than built them. What I bring is the control view: where the boundaries are and what has to be enforced at each.'	'Where in this design is scope enforced — in the application or in a service beneath it?'
'This is just a BI problem.'	'The disciplines are the same — catalogue, master data, reconciliation, lineage, access control. The consequences are not: the requirements come from a legal instrument, and over-delivery is a breach. That changes what we optimise for.'	'Which of the retained fields were chosen because the feed offered them rather than because the notice named them?'
'Our MDM engine will handle identity resolution.'	'Good — is its matching deterministic? Probabilistic matching is standard in MDM and it is ruled out here because a match score cannot be explained in a witness box.'	'Can you show me the resolution for a number that was reassigned mid-period, with the gap?'
'We have comprehensive logging.'	'Logging is the raw material; what I need is a reconstructable record per disclosure of what was requested, what was authorised and what was returned — and it has to outlive the data.'	'For a disclosure from six months ago, show me the three field sets.'
'The data quality framework covers it.'	'Which dimension does it measure completeness against? If it is the feed's own count, a stopped feed scores 100%.'	'What was yesterday's expected versus actual volume for the NAT feed?'
'Deletion is handled by the lifecycle policy.'	'For the live store, yes. What about the replicas, the correlated copy, the indexes and the backups?'	'When a partition expires, is its key destroyed?'
'We'll use the charging records; they're already consolidated.'	'As a first release, fine. They reflect what is billable, they change with the tariff, and they omit unanswered calls. I want a field-by-field gap analysis against the notice and a named consultee on billing changes.'	'Who is told before a tariff change goes live, and is the capability-impact question on the change record?'
'The vendor platform does all of this out of the box.'	'Then the policy — classification, scope rules, correlation, nil-return vocabulary — lives inside their product. I want it in configuration we own and can move.'	'If we changed supplier, what would we have to rebuild?'

TALKING POINT

The pattern in every script: acknowledge the parallel, state the difference in one sentence, and ask a question that can only be answered by showing you something. You do not need to have built the warehouse to ask to see the reconciliation report.